

Indexing(B-tree,R-tree)

Lee,Jae-Yong
October 10,2002

1

Presentation Outline

- Motivation
- B tree
- B+ tree, B* tree
- R tree
- R* tree
- SS tree,SR tree

2

Why indexing?

- When data is too large to fit in memory, then the number of disk accesses is important.
- A disk access is expensive and worth about 200,000 instructions.
- To minimize disk accesses


3

m-way search tree

- A generalization of binary search tree
- Each node has at most m child-pointers
- If $k \leq m$ is the number of children, then the node has exactly $k-1$ keys.
- The tree is ordered in keys.

4

Multiway Search Trees (cont.)



ITB421-1998/2

© Frederic Marx, QUT

5

5


B tree(Definition)

- B-tree of order $2d$ is an $(2d+1)$ -way search tree such that:
 - all leaf nodes are on the same level.
 - completely balanced.
 - All non-leaf nodes(except root node) have at most $2d$ keys and $2d+1$ pointers and at least d keys and $d+1$ pointers.
 - The root has at least one key and 2 child pointers.

6

6

Example(B-tree)



7

Height of B-tree

- What is the maximum height of a B-tree with n keys?
- The maximum height of a B-tree will give an upper bound on the number of disk access.
- The minimum number of keys in a B-tree of order $2d$ and depth h :

$$1 + 2d + 2d(d+1) + 2d(d+1)^2 + \dots + 2d(d+1)^{h-1}$$

$$= 1 + \sum_{i=1}^h 2d(d+1)^{i-1}$$

- The maximum height of a B-tree with n keys:

$$\log_{d+1} \frac{n}{2d} = O(\log_d n)$$


8

Search(Find)

- B-tree is always completely balanced.
- The length of retrieval path is at most the height of the B-tree.
- The cost of a search is $O(\log_d n)$.

9

Example(Find 70)



10

Insert

- First, a find proceeds from the root to locate the proper leaf node for insertion.
 $\Rightarrow O(\log_d n)$
- If the leaf node is already full, split it.
 Otherwise, done.

11

Splitting a node

- $2d+1$ keys are split to 2 new nodes.
- One node has the smallest d keys and the other has the largest d keys.
- A separator key is promoted to the parent node.
- If the parent node is already full, the parent node is split too.


12

Splitting a node(cont.)

- In the worst case, splitting propagates to the root node and the tree increases in the height by one level.


13

Example(Splitting, Insert 67)



14

Example(cont.)



15

Insert cost

- Find operation : $O(\log_d n)$
- Splitting : $O(\log_d n)$
- Total cost is at most doubled as a find .
- Usually, nodes in the path reside in memory after the find operation.

16

Delete

- First, a find operation to locate the node
- The key can be found in either a leaf or a nonleaf
- In Insert, the key is always inserted to a leaf node
- Nonleaf case
 - Find an adjacent key in a subtree of the deleted key (found in a left-most leaf node of the right subtree)
 - Move it to the deleted position
 - Process in the leaf node


17

Underflow during delete

- Underflow
 - Less than d keys occupy the node($d-1$ keys)
 - Need to redistribute the keys to a neighbor node(including their parent key)
- Redistribution
 - Evenly divide keys between two neighbors
 - Need at least $2d+1$ ($d+d+1$) keys to redistribute
 - What if the neighbor node does not have enough keys (has just d nodes) \rightarrow concatenation

18

Redistribution(Delete 6)




19

Underflow during delete(cont.)

- Concatenation
 - Inverse of Splitting
 - Keys in two neighbor nodes and a separator key in the parent node
 - The merged node has 2d keys(full).
 - The parent node may be concatenated.
 - May be propagated all the way to the root node
 - Decreasing the height of the B-tree

20

Concatenation(delete 3)



21

Delete cost

- Same as Insert operation

22

Running time

- Running time of each operation is directly related to the height of the tree
- The height of a B-tree depends on :
 - the number of keys in the tree
 - the branching factor($2d+1$)
 - Practical limit on a size of node
 - Usually a block size(page size)
 - the storage utilization
 - B* tree(at least 2/3 of a node is full)

23

Sequential processing in a B-tree

- Preorder tree walk
- Stacks the nodes along a path from the root
- Require space for at least the height
- B-tree does not do well in a sequential processing
- B+ tree


24

B+ tree

- Independent index and sequence sets
- Sequence set(leaves)
 - all actual keys reside in the leaves.
 - leaf nodes are linked for a sequential processing
- Index set(non-leaves)
 - Just a roadmap to reach a leaf node
 - Not necessarily actual keys
- Better sequential processing

25

Example(B+ tree)



26

B* tree

- Each node is at least 2/3 full.
- Increasing storage utilization
 - speeding up search operations
- Split
 - Redistribution
 - Delay the split until two siblings are full
 - Reduce the number of splits
 - 2 nodes are divided into 3 nodes

27

R-Tree(Motivation)

- Classical indexing (like B-Trees) are not suitable for multidimensional spatial searching

28

Spatial Objects

- Data objects of non-zero size in multi-dimensional spaces
 - ex) a country in a map
- Queries are also spatial
 - ex) Find all objects in an area

29

Nodes in R-tree

- Leaf node
 - has entries of the form (I, tuple-identifier)
 - tuple-identifier: a reference to an object in DB
 - $I = (I_0, I_1, \dots, I_{n-1})$: an n-dimensional bounding box
- Nonleaf node
 - has entries of the form (I, child-pointer)
 - child-pointer
 - I : covers all rectangles in the lower node's entries

30

R-tree

- M is the max number of entries in a node.
- $m \leq \frac{M}{2}$ is the min number of entries
- m is not fixed and varies for performance tuning.



31

R-tree(Definition)

- Every node contains between m and M entries.(except the root)
- The root has at least two children.
- All leaves appear on the same level
- completely balanced
- I is a rectangle that contains all objects in lower levels


32

Example




33

An Example



An Example



R-tree

- Nonleaf nodes serve as a set of indices.
 - B+ tree
- Bounding boxes in a same level are not disjoint.
 - A level in B-tree consists of disjoint intervals.
- Each entry in a node represents an area in space.

36

Search

- Bounding boxes in each level may be overlapped.
 - More than one subtrees under a node may need to be searched
 - No good worst case performance
 - Need some techniques to minimize searches
 - Reduce dead-space
 - Update algorithms maintains “good” trees.

37


Search(cont.)

- Query : Search rectangle S
- Search all children if their bounding boxes overlap with S.
- All leaf nodes have a tuple identifier
- The cost depends on the structure of the tree
 - Need to minimize the area of MBR of nodes

38

Example(Search)

A 2-dim Example (R-tree)



39

Insert

- Like the B-tree, sometimes it is necessary to split a node
- Since the decision whether to visit a node depends on whether its covering rectangle overlaps the search area, areas of two new nodes should be minimized
 - All approaches of optimizing the retrieval performance have to be applied during Insert operation

40

Insert(cont.)

- Similar to B-tree in handling split and simple insertions
- ChooseLeaf: selects a leaf node L in which to place the new entry
- The algorithm tries to find the region whose rectangle needs least enlargement to include the new entry

41

Delete

- Like in B-Tree
 - But find a leaf to be deleted (like B+ tree)
- Sometimes you need to condense a tree due to the elimination of one node


42

Delete(Condense)

- Underflow in the node after a deletion
- An underflow may be propagated.
- If underflow in the node, N,
 - Delete the index entry of parent node, P
 - Add N to the set of eliminated nodes, Q
 - When the propagation is done, reinsert all entries of nodes in Q
 - Need to reinsert a proper level in the tree

43

Node Splitting



04/14/1992

CMSC 438 - Spring '92

32 of 40

Node splitting algorithm

- “Good” split can reduce the number of nodes during subsequent searches
 - Exhaustive algorithm
 - Quadratic algorithm(Heuristic)
 - Linear algorithm(Heuristic)

45

Exhaustive method

- Generate all possible groupings (2^M possibilities)
- M is typically 50
- Too expensive

46

Quadratic method

- picks 2 of the M+1 entries to be seeds of the 2 new groups -> $O(dM^2)$
 - the pair that would waste the most are if both were put in the same group
- then assign to groups one entry at a time
- at each step the area expansion required to add each remaining entry to group is calculated, and the entry assigned is the one showing the greatest difference between the 2 groups -> $O(dM^2)$

47

Linear method

- Different seed picking -> $O(dM)$
 - find the entry whose rectangle has the highest low side, and the one with the lowest high side. Record the separations for each dimension ->
 - normalize the separations by dividing by the width of the entire set
 - choose the pair with the greatest separation along any dimension.
- Assign each entry to a group -> $O(dM)$
 - Pick any entry and then assign the group which requires less expansion

48

Performance Test

- evaluate m (minimum number of entries) as a function of M (maximum number) and the heuristics
- The linear node-split algorithm proved to be as good as more expensive techniques
 - faster and the slightly worse quality of the splits did not affect search performance noticeably

49

Problems in R-tree

- Just try to minimize the area of each enclosing rectangle in the index nodes.
- What can we consider to produce “good” trees?
 - R* tree
- What if the query is not a rectangle?
 - SS tree, SR tree

50

R* tree

- Optimization Criteria:
 - (O1) Area covered by an index MBR
 - (O2) Overlap between directory MBR
 - (O3) Margin of a directory rectangle
 - Next slide.
 - (O4) Storage utilization
- Sometimes it is impossible to optimize all the above criteria at the same time

51

Overlap & Margin

- (O2) Overlap between directory MBR
 - Sum of overlapped areas of all pairs of entries in the MBR
- (O3) Margin of a directory rectangle
 - Margin : sum of lengths of the edges in MBR
 - Minimizing the margin of Rectangle
 - Rectangle can be more quadratic. Quadratic object can be packed easier. \rightarrow smaller MBR

52

R*-tree(cont.)

- ChooseSubtree
 - Leaf node:
 - Least overlap enlargement
 - Least area enlargement(R-tree)
 - Smaller area
 - Nonleaf node
 - Least area enlargement(R-tree)
 - Smaller area

53

Split

- SplitNode
 - Choose the axis to split
 - Choose 2 groups along the chosen axis
- ChooseSplitAxis
 - Along each axis, sort rectangles and break them into 2 groups($M-2m+2$ ways)
 - Compute the sum S of all margin-values of each pair of groups. Choose the one that minimizes S
- ChooseSplitIndex
 - Along the chosen axis, choose the grouping that gives the minimum overlap-value

54

Forced Reinsert

- The structure of tree is non-deterministic
- Dynamic reorganization.
- Defer splits, by forced-reinsert
 - Instead splitting, temporarily delete some entries, shrink overflowing MBR, and reinsert those entries
- Which ones to re-insert?(next slide)
- How many? 30% (by experiments)

55

Forced Reinsert(Details)

- Calculate centers of MBRs of all entries in the overflowed node and the center of MBR of the node
- Remove p farthest entries from the node
- Then reinsert the entries
- If at least one entry is inserted into different node, no splitting the overflowed node.
 - But may cause other nodes to be overflowed

56

SS & SR-tree

- An index structure for high-dimensional nearest neighbor queries
- the similarity queries on feature vectors which is widely used for performing content-based retrieval of images

57